

Skeleton-Aware Graph Embeddings for Multi-Person Pose Grouping

by

Dean Spencer Mills

Submitted in partial fulfillment of the requirements for the degree
Master of Engineering (Computer Engineering)

in the

Department of Electrical, Electronic and Computer Engineering
Faculty of Engineering, Built Environment and Information Technology

UNIVERSITY OF PRETORIA

January 2025

SUMMARY

Skeleton-Aware Graph Embeddings for Multi-Person Pose Grouping

by

Dean Spencer Mills

Supervisor: <to be completed>
Co-supervisor: <to be completed> [optional line]
Department: Electrical, Electronic and Computer Engineering
University: University of Pretoria
Degree: Master of Engineering (Computer Engineering)
Keywords: Keyword, keyword, keyword, keyword, keyword, keyword, keyword,
keyword, keyword, keyword, keyword, keyword

This is the EECE dissertation/thesis template for Latex, version 2025-11-24. Summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one.

Summary paragraph two summary paragraph two summary paragraph two summary paragraph two
summary paragraph two summary paragraph two summary paragraph two summary paragraph two
summary paragraph two summary paragraph two summary paragraph two summary paragraph two
summary paragraph two summary paragraph two summary paragraph two summary paragraph two

summary paragraph two summary paragraph two summary paragraph two summary paragraph two
summary paragraph two summary paragraph two.

Summary paragraph three summary paragraph three summary paragraph three summary paragraph
three summary paragraph three summary paragraph three summary paragraph three summary paragraph
three summary paragraph three summary paragraph three summary paragraph three summary paragraph
three summary paragraph three summary paragraph three.

LIST OF ABBREVIATIONS

ANN	artificial neural network
EMI	electromagnetic interference
MCMC	Markov chain Monte Carlo
SVM	support vector machine

Note: Make sure the list of abbreviations is sorted alphabetically. This list is formatted as a table with no borders. Insert rows as needed.

LIST OF SYMBOLS [OPTIONAL]

I	current
R	resistance
V	voltage

Note: Make sure the list of symbols is sorted alphabetically, and/or according to the Greek alphabet. The list is formatted as a table with no borders. Insert rows as needed.

TABLE OF CONTENTS

CHAPTER 1	INTRODUCTION	1
1.1	CONTEXT	1
1.2	PROBLEM STATEMENT	1
1.3	RESEARCH GAP	1
1.4	RESEARCH OBJECTIVE AND RESEARCH QUESTIONS	1
1.5	APPROACH	2
1.6	RESEARCH CONTRIBUTIONS	3
1.7	RESEARCH OUTPUTS	3
1.8	OVERVIEW OF STUDY	3
1.9	DECLARATION REGARDING USE OF GENERATIVE AI	4
CHAPTER 2	LITERATURE REVIEW	5
2.1	MULTI-PERSON HUMAN POSE ESTIMATION	5
2.1.1	Top-Down Approaches	6
2.1.2	Bottom-Up Approaches	6
2.1.3	Transformer-Based Approaches	6
2.1.4	Heatmap and Regression Trends	7
2.1.5	Foundation Models for Pose Estimation	8
2.1.6	Benchmarks	8
2.1.7	Pose Tracking	8
2.1.8	Position of This Dissertation	9
2.2	GROUPING IN BOTTOM-UP POSE ESTIMATION	9
2.3	GRAPH NEURAL NETWORKS FOR POSE	9
2.4	REPRESENTATION LEARNING FOR CLUSTERING	9
2.5	GEOMETRIC EMBEDDING SPACES	9

2.6	RESEARCH GAP	10
CHAPTER 3	METHODOLOGY	11
3.1	PRELIMINARIES — GRAPH NEURAL NETWORKS	11
3.1.1	The Message-Passing Framework	11
3.1.2	Graph Convolutional Networks (GCN)	12
3.1.3	Graph Attention Networks (GAT and GATv2)	13
3.2	PRELIMINARIES — CLUSTERING AND LEARNED GROUPING	14
3.2.1	k -means and constrained k -means	14
3.2.2	Deep embedded clustering (DEC)	14
3.2.3	Slot attention	14
3.2.4	Graph partitioning (MinCut, NormCut)	14
3.2.5	Deep modularity networks (DMoN)	14
3.3	PROBLEM FORMULATION	14
3.4	DATASET PIPELINE	15
3.5	BASELINE: STANDARD GAT + CONTRASTIVE + KNN	15
3.6	LEARNED GROUPING HEADS	15
3.6.1	DEC applied to pose grouping	15
3.6.2	Slot attention applied to pose grouping	15
3.6.3	Graph partitioning applied to pose grouping	15
3.7	SA-DMON — SKELETON-AWARE DEEP MODULARITY NETWORKS	15
3.8	PG-GAT — POSE GROUPING GRAPH ATTENTION NETWORK	15
3.8.1	Design intuition	15
3.8.2	Type-pair attention modulation	16
3.8.3	Skeleton-aware positional encoding	16
3.8.4	Type-repulsion loss term	16
3.8.5	Training procedure	16
3.9	HYPERBOLIC PG-GAT	16
3.10	TRIGAT — TRIPLET-BASED PRIMITIVE	16
3.11	INTEGRATION WITH A REALISTIC DETECTOR	16
3.12	EVALUATION PROTOCOL AND METRICS	16
CHAPTER 4	EXPERIMENTAL RESULTS	18
4.1	BASELINE: STANDARD GAT	18

4.2	FAILED LEARNED GROUPING HEADS	18
4.2.1	DEC	18
4.2.2	Slot attention	18
4.2.3	Graph partitioning	18
4.2.4	Cross-cutting finding	19
4.3	SA-DMON — TEN EXPERIMENTS	19
4.4	PG-GAT ABLATIONS	19
4.5	COCO FINE-TUNING	19
4.6	HYPERBOLIC PG-GAT	19
4.7	TRIGAT — NEGATIVE RESULT	19
4.8	END-TO-END INTEGRATION WITH HIGHERHRNET	20
4.9	MODEL SIZE AND MODULARITY	20
4.10	VISUAL-FEATURE AUGMENTATION (PG-GAT V2) — NEGATIVE RESULT	20
CHAPTER 5	DISCUSSION	21
5.1	WHY THE LEARNED GROUPING HEADS FAILED	21
5.2	WHY SA-DMON FAILED	21
5.3	WHY PG-GAT WORKS	21
5.4	INTERPRETING THE HYPERBOLIC RESULT	22
5.5	WHY TRIGAT UNDERPERFORMED	22
5.6	COMPARISON WITH HIGHERHRNET AE — THE MODULARITY TRADE-OFF	22
5.7	ANSWERS TO THE RESEARCH QUESTIONS	22
5.8	LIMITATIONS AND THREATS TO VALIDITY	22
CHAPTER 6	CONCLUSION	23
6.1	SUMMARY OF CONTRIBUTIONS	23
6.2	ANSWERS TO THE RESEARCH QUESTIONS	23
6.3	FUTURE WORK	23
REFERENCES		24
APPENDIX A	DERIVATION OF HIGHER ORDER MODELS	30
A.1	EXAMPLE HEADING	30
A.1.1	Subheading example	30

LIST OF FIGURES

[THE LIST OF FIGURES IS OPTIONAL]

LIST OF TABLES

[THE LIST OF TABLES IS OPTIONAL]

CHAPTER 1 INTRODUCTION

This chapter motivates multi-person pose grouping as a separable problem within bottom-up human pose estimation (HPE), identifies the gap in modular detector-agnostic grouping methods, and sets out the research questions, approach, and contributions that the rest of the dissertation defends.

1.1 CONTEXT

Introduces the detection/grouping split in bottom-up HPE: modern detectors (HigherHRNet, HRNet) produce accurate per-joint heatmaps; assigning those joints to person identities is the remaining bottleneck. Summarises why joint-to-person assignment is intrinsically difficult (unknown K , type constraints, occlusion, scale) and why it deserves to be studied as a module in its own right.

1.2 PROBLEM STATEMENT

Formulates the task: given N detected keypoints with joint-type labels $t_i \in \{1, \dots, 17\}$ from an image containing an unknown number of people K , assign each keypoint to one of K person groups under the constraint that every person has at most one keypoint of each type.

1.3 RESEARCH GAP

Existing grouping methods (associative embeddings in HigherHRNet, part-affinity fields in OpenPose, HGG, CenterGroup) are tightly coupled to their host detectors and trained jointly with them. A *modular* grouping module that operates on any detector’s keypoint output is largely unstudied. This is the gap this work addresses.

1.4 RESEARCH OBJECTIVE AND RESEARCH QUESTIONS

The objective is to design and evaluate a detector-independent, GNN-based pose grouping module. The investigation addresses the following research questions. Bracketed short answers anticipate the findings from Chapter 4; full support is provided there.

- **RQ1 — The lever.** Is the bottleneck for clustering-based pose grouping the embedding, or the grouping algorithm? *[Answer: the embedding. Four learned grouping heads (DEC, slot attention, graph partitioning, SA-DMoN) fail to beat plain kNN on the same embeddings; strengthening the embedding (PG-GAT) then beats every prior configuration.]*
- **RQ2 — Inductive biases.** Which skeleton-aware modifications to a GATv2 embedding improve pose grouping, and by how much? *[Answer: three modifications — type-pair attention modulation, skeleton-aware positional encoding, and a type-repulsion loss term — each contribute independently; the full combination is the best configuration.]*
- **RQ3 — Detector independence.** Can a grouping module trained separately from the detector match an integrated end-to-end grouping mechanism (HigherHRNet’s associative embedding head) on a real-image benchmark? *[Answer: not fully. PG-GAT + kNN end-to-end is approximately 0.04 PGA below HigherHRNet AE on COCO, with detector-independence as the trade-off.]*
- **RQ4 — Embedding geometry.** Does the choice of embedding space (Euclidean vs hyperbolic) affect pose grouping quality? *[Answer: yes before COCO fine-tuning, where hyperbolic at $d = 32$ matches Euclidean at $d = 128$ (an efficiency edge). After fine-tuning the advantage disappears — reported as an efficiency trade, not a quality win.]*
- **RQ5 — Graph primitive.** Is a higher-order graph primitive (skeleton triplets) a better node for pose grouping than single joints? *[Answer: no. TriGAT underperforms joint-level PG-GAT on every comparison, with the voting step (triplet clusters to joint labels) as the dominant failure mode.]*
- **RQ6 — Fine-tuning.** How much does COCO fine-tuning lift a synthetically-trained embedding, and does it damage performance on the synthetic domain? *[Answer: +7.5 PGA on COCO (kNN), the single largest gain of any experiment in this work; synthetic performance is not damaged ($0.910 \rightarrow 0.928$).]*

1.5 APPROACH

Train a graph attention network with contrastive loss on detected keypoints; cluster at test time with kNN. The iteration order mirrors the investigation: standard GAT baseline $\rightarrow$ a series of learned grouping heads (DEC, slot attention, graph partitioning), all of which failed to beat kNN $\rightarrow$ SA-DMoN, a skeleton-aware modularity objective, which also failed $\rightarrow$ PG-GAT, which keeps the clustering head simple but strengthens the embedding with three skeleton-aware modifications. Evaluate on synthetic

data (controlled K , perfect GT) and COCO val2017 (zero-shot and fine-tuned). Close the loop by integrating with HigherHRNet for end-to-end validation.

1.6 RESEARCH CONTRIBUTIONS

The contributions of this work are:

1. **PG-GAT**, a skeleton-aware graph attention network for pose-keypoint embedding. Three modifications over standard GATv2: type-pair attention modulation, skeleton-aware positional encoding, and a type-repulsion term in the contrastive loss.
2. **Empirical evidence that the embedding, not the grouping head, is the decisive lever** for clustering-based pose grouping. Documented across four learned-head approaches (DEC, slot attention, graph partitioning, and ten SA-DMoN variants), none of which outperformed plain kNN on the same embeddings.
3. **A COCO fine-tuning protocol** that lifts PG-GAT to 0.957 PGA on COCO val2017 with kNN clustering, and end-to-end to 0.942 when paired with HigherHRNet detection.
4. **An alternative-embedding-space study** (hyperbolic PG-GAT in the Lorentz model): hyperbolic at $d = 32$ matches Euclidean at $d = 128$ pre-fine-tune, but the advantage disappears after COCO fine-tuning — documented as an efficiency trade, not a quality win.
5. **An alternative-primitive study** (TriGAT, triplet nodes): a negative result reported with the failure-mode analysis (voting step propagates errors).

1.7 RESEARCH OUTPUTS

The primary research output is a journal paper on PG-GAT, targeted at *Image and Vision Computing* (Q1, IF 5.34), with *Pattern Recognition* (Q1, IF 9.84) as a stretch target.

1.8 OVERVIEW OF STUDY

Chapter 2 reviews the literature on bottom-up human pose estimation, graph neural networks, and clustering methods relevant to pose grouping. Chapter 3 presents the mathematical preliminaries (GNN, GCN, GAT/GATv2, DEC, slot attention, graph partitioning, DMoN) and describes the methodology behind each method investigated, culminating in PG-GAT. Chapter 4 reports experimental results across the methods. Chapter 5 discusses what the results reveal about the problem. Chapter 6 concludes and indicates future work.

1.9 DECLARATION REGARDING USE OF GENERATIVE AI

Please declare if and how generative AI was used in this work. This must be done following the steps (adapted for dissertations/theses) outlined in 8.2.1.B.10 in the IEEE Publication Services and Products Board Operations Manual¹:

The use of content generated by artificial intelligence (AI) in ... [a dissertation/thesis] (including but not limited to text, figures, images, and code) shall be disclosed ... The AI system used shall be identified, and specific sections of the ... [dissertation/thesis] that use AI-generated content shall be identified and accompanied by a brief explanation regarding the level at which the AI system was used to generate the content.

While the use of AI systems for editing and grammar enhancement is common practice, its use in this role must also be disclosed.

¹<https://pspb.ieee.org/images/files/PSPB/opsmanual.pdf>

CHAPTER 2 LITERATURE REVIEW

This chapter situates multi-person pose grouping within the broader human pose estimation literature, reviews the existing families of grouping mechanisms, and surveys the graph neural network and representation-learning work that PG-GAT builds on. The review closes by identifying the specific gap that this dissertation addresses.

2.1 MULTI-PERSON HUMAN POSE ESTIMATION

Human pose estimation (HPE) is the problem of localising a fixed set of anatomical keypoints — such as the seventeen joints of the COCO skeleton [1] — on every person present in an image. The dissertation restricts attention to the 2D, multi-person, image-level formulation: per-image keypoint coordinates for an unknown number of people, without temporal context and without 3D reconstruction. Three surveys provide broader scope and historical depth: Chen et al. [2] cover monocular HPE end-to-end, Zheng et al. [3] provide a more recent deep-learning-focused survey, and Andriluka et al. [4] anchor the modern benchmark culture with the MPII dataset.

The convolutional era of HPE opens with Toshev and Szegedy’s DeepPose [5], which framed keypoint localisation as direct coordinate regression from a deep network. Subsequent architectures shifted to dense heatmap prediction: convolutional pose machines [6] introduced iterative belief-map refinement, and the stacked hourglass network [7] established the encoder–decoder-with-skip-connections topology that underpins most modern single-person estimators. Xiao et al.’s “simple baselines” work [8] showed that a ResNet trunk with three deconvolution layers matched hourglass accuracy at lower complexity, establishing a strong baseline that subsequent methods are still measured against. The HRNet family [9] pushed this further with parallel multi-resolution branches that retain high-resolution features throughout the network, which became the de facto backbone for most top-down systems and for the bottom-up HigherHRNet variant [10] discussed in Section 2.2.

2.1.1 Top-Down Approaches

Top-down systems decompose multi-person HPE into two stages: a person detector produces bounding boxes, and a single-person estimator runs on each crop. The canonical instantiations of this architecture are Mask R-CNN [11], which attaches a per-keypoint mask head to the Faster R-CNN detector, and AlphaPose [12], which couples a person-detection front-end with a symmetric spatial transformer and a parametric pose non-maximum suppression step. Top-down methods inherit the strong per-person localisation accuracy of the underlying single-person estimator and consequently dominate the COCO leaderboards across most years. Two structural limitations apply. First, runtime scales linearly with the number of detected people, which becomes prohibitive in dense scenes. Second, performance is bounded by detection quality: missed or over-merged person detections propagate directly to missed or incorrectly grouped keypoints, with no opportunity for the pose estimator to recover.

2.1.2 Bottom-Up Approaches

Bottom-up methods reverse the dependency: all keypoints in the image are detected first, and a separate grouping mechanism assigns each keypoint to a person. Runtime is then approximately constant in the number of people, and detection errors are localised to individual keypoints rather than to entire person instances. Three distinct grouping lineages have emerged within the bottom-up family, and the literature on the grouping mechanisms themselves is reviewed in Section 2.2; the present section names them only to map the bottom-up landscape.

The first lineage, originating with OpenPose [13] and extended by PifPaf [14], encodes pairwise keypoint associations as part affinity fields and recovers person assignments by integrating along inferred limb directions. The second lineage, introduced by Newell et al. as associative embedding [15] and extended to high-resolution backbones by HigherHRNet [10], attaches a learned tag vector to each keypoint and groups by tag-vector similarity at inference time. The third lineage, exemplified by DEKR [16] and CenterGroup [17], regresses keypoint offsets from disentangled person centres and assembles full skeletons by following the offset vectors. End-to-end transformer methods (Section 2.1.3) collapse the detect-and-group structure entirely; they are catalogued separately here because they no longer admit a modular grouping stage.

2.1.3 Transformer-Based Approaches

The transformer architecture has been applied to HPE both as a backbone replacement and as a complete end-to-end formulation. TokenPose [18] demonstrated that a pure-transformer network treating each keypoint as a learned token could match the heatmap accuracy of CNN baselines at comparable

parameter counts. HRFormer [19] preserved the parallel-resolution structure of HRNet while replacing convolutions with local-window self-attention, recovering high-resolution feature representations in the transformer family. ViTPose [20] took the opposite design choice — a plain monolithic Vision Transformer as backbone — and showed that this minimal architecture scales cleanly to model sizes from 100M up to 1B parameters. Its successor ViTPose++ [21] adds a mixture-of-experts formulation to cover the wholebody, hand, and animal-pose distributions within a single trunk.

A separate transformer line of work reformulates multi-person HPE as set prediction, dispensing with the explicit grouping stage altogether. PETR [22] was the first end-to-end pose detector to apply the DETR set-prediction paradigm to human keypoints, predicting a fixed number of pose queries with bipartite matching at training time. QueryPose [23] refined this with part-level queries to sparsify the attention pattern. ED-Pose [24] and GroupPose [25] further developed the query-based formulation, reaching competitive performance on the COCO benchmark without any post-hoc grouping step. These methods couple grouping tightly to the backbone: they do not produce explicit per-keypoint embeddings that a downstream module could group separately, and the grouping mechanism is consequently not transferable to checkpoints from other detectors.

2.1.4 Heatmap and Regression Trends

The traditional bottleneck of regression-based pose estimation is the discontinuity of the argmax operator that maps a continuous output to a discrete keypoint coordinate. Sun et al.’s integral pose regression [26] bridged the heatmap and regression paradigms with a differentiable soft-argmax, and DSNT [27] formalised the same idea as a normalised spatial-coordinate layer. A direct probabilistic treatment was given by residual log-likelihood estimation (RLE) [28], which models the residual between predicted and ground-truth coordinates with a flow-based density and was the first regression method to exceed heatmap accuracy on multi-person COCO. SimCC [29] reformulated coordinate prediction as a two-axis classification problem, producing the underlying mechanism for the real-time RTMPose family [30]. RTMO [31] unifies detection and regression in a single stage, and RTMW [32] extends the family to whole-body pose at real-time rates. These methods are relevant to this dissertation only insofar as they produce per-keypoint coordinates; the grouping module proposed in Chapter 3 operates on detected coordinates and is agnostic to whether they originated from a heatmap or a regressor.

2.1.5 Foundation Models for Pose Estimation

Recent work has pushed the backbone trajectory toward foundation-model scale. Sapiens [33] pretrains Vision Transformer backbones from 0.3 to 2 billion parameters on 300M curated human images, producing a unified human-vision backbone that reaches state-of-the-art accuracy across pose estimation, segmentation, depth, and surface normal prediction. The same trend toward promptable, dataset-unified methods is represented by X-Pose [34], which detects arbitrary keypoint categories conditioned on a textual or visual prompt across thirteen datasets, and by the compositional-token approach of PCT [35], which represents poses as discrete codebook tokens to improve occlusion robustness. The foundation-model trajectory matters for this dissertation as context rather than as comparison: a sufficiently large pretrained backbone reduces the per-keypoint localisation error to near saturation, which sharpens the case for studying grouping as a modular problem downstream of detection.

2.1.6 Benchmarks

Three datasets dominate the 2D multi-person literature. MPII Human Pose [4] established the modern benchmark protocol with 25K images and approximately 40K annotated people across 410 activities; its focus is single-person crops and it has been largely superseded for multi-person research. COCO [1] is the de facto multi-person benchmark, with seventeen keypoints per person and approximately 250K annotated instances across the train, val, and test partitions; val2017 is the standard reporting set and the one used for the COCO evaluations in Chapter 4. CrowdPose [36] targets the dense-crowd regime explicitly, with a crowd-index metric and approximately 20K images selected for inter-person occlusion; it is the appropriate benchmark for stress-testing grouping methods, although it is not used in this dissertation. The OCHuman benchmark, introduced alongside Pose2Seg [37], isolates occluded humans and is similarly relevant to grouping robustness.

2.1.7 Pose Tracking

Multi-person 2D pose tracking extends per-image HPE with temporal association across frames. PoseFlow [38] is the canonical online tracker that propagates per-person identities through pose-similarity matching in a small temporal window. LightTrack [39] provides a top-down tracking framework with a learned re-identification head. The PoseTrack benchmark [40] is the standard evaluation set for this video-pose formulation. Pose tracking is out of scope for this dissertation, which is restricted to per-image grouping, but the per-image grouping module developed in Chapter 3 is compatible with downstream temporal association via standard ReID techniques.

2.1.8 Position of This Dissertation

The HPE literature has converged on a small set of strong single-image detection backbones whose per-keypoint localisation accuracy is approaching saturation. In bottom-up systems, three established grouping lineages (part-affinity fields, associative embedding, centre-regression) all either tie grouping to a specific backbone or fold it into the detector training objective, preventing detector substitution after training. End-to-end transformer pose detectors collapse detection and grouping into a single set-prediction pass, with the same coupling consequence. The lineage of modular, detector-agnostic grouping methods that operate on already-detected keypoints is consequently sparser, and this dissertation contributes to that lineage. The next two sections develop the grouping mechanism (Section 2.2) and the graph neural network machinery (Section 2.3) on which the contributed method (PG-GAT, Chapter 3) is built.

2.2 GROUPING IN BOTTOM-UP POSE ESTIMATION

Review of the major grouping mechanisms in the bottom-up literature: part-affinity fields (OpenPose), associative embeddings (Newell et al.), PersonLab mid-range offsets, HGG (graph-based), CenterGroup. How each binds tightly to its host detector and why that limits modularity.

2.3 GRAPH NEURAL NETWORKS FOR POSE

Existing graph-based approaches in the pose literature: STGAT (skeleton-based action recognition), GAST-Net (single-person 3D pose from video), skeleton-aware GCN variants. How these differ from the multi-person keypoint grouping task — none of the existing “skeleton-aware” networks target grouping, which motivates both the PG-GAT contribution and the naming choice (avoiding collision with STGAT/GAST-Net descriptors).

2.4 REPRESENTATION LEARNING FOR CLUSTERING

Contrastive learning (InfoNCE, SimCLR), triplet losses, metric learning. Deep clustering lineage: DEC (Xie et al.), IDEC, DeepCluster, DMoN. What these methods assume about the feature space and how those assumptions interact with the pose-grouping task.

2.5 GEOMETRIC EMBEDDING SPACES

Euclidean versus non-Euclidean embedding spaces. Hyperbolic embeddings (Poincaré ball, Lorentz model) and their suitability for tree-structured data. Relevance to skeletons, which are trees. Sets up the motivation for the hyperbolic PG-GAT variant evaluated in Chapter 4.

2.6 RESEARCH GAP

Synthesises the review: modular grouping modules for bottom-up HPE are under-explored; existing graph-based pose networks target different tasks; the clustering-head literature has not been systematically applied to pose grouping; and the choice of embedding space (Euclidean vs hyperbolic) is untested in this domain. PG-GAT and its ablations address these gaps.

CHAPTER 3 METHODOLOGY

This chapter presents the mathematical preliminaries that every method in this dissertation builds on, the problem formulation and data pipeline, and the methodology for each approach investigated — from the standard GAT baseline, through the series of learned grouping heads that were tried and failed, through SA-DMoN, and finally to PG-GAT. Two variants (hyperbolic PG-GAT, TriGAT) and the integration with a realistic detector are described last.

3.1 PRELIMINARIES — GRAPH NEURAL NETWORKS

A human skeleton is a graph: seventeen keypoints form the nodes and the skeletal connections form the edges. A graph neural network (GNN) embeds each node while conditioning on the structure of the graph — the natural choice for producing pose-grouping embeddings that are aware of joint-type and connectivity information. This section introduces the message-passing framework that unifies modern GNN variants, derives the graph convolutional network (GCN) as the canonical spectral instance, and develops the graph attention network (GAT), culminating in the GATv2 correction that is used as the attention primitive in every experiment in this dissertation.

3.1.1 The Message-Passing Framework

Let $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ denote an undirected graph with N nodes and edge set $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$. Each node $i \in \mathcal{V}$ is associated with a feature vector $\mathbf{h}_i^{(0)} \in \mathbb{R}^{d_0}$, and each edge $(i, j) \in \mathcal{E}$ optionally carries an edge feature $\mathbf{e}_{ij}$. The neighbourhood of node i is $\mathcal{N}(i) = \{j : (i, j) \in \mathcal{E}\}$.

Almost every modern GNN variant can be expressed as iterated message passing [41]. At layer $l + 1$, each node receives a message from each of its neighbours, aggregates them with a permutation-invariant operator, and updates its representation accordingly. Formally,

$$\mathbf{m}_{i \leftarrow j}^{(l+1)} = \text{MSG}^{(l)}\left(\mathbf{h}_i^{(l)}, \mathbf{h}_j^{(l)}, \mathbf{e}_{ij}\right), \quad (3.1)$$

$$\mathbf{h}_i^{(l+1)} = \text{UPDATE}^{(l)} \left(\mathbf{h}_i^{(l)}, \bigoplus_{j \in \mathcal{N}(i)} \mathbf{m}_{i \leftarrow j}^{(l+1)} \right), \quad (3.2)$$

where $\bigoplus$ denotes a permutation-invariant aggregator (sum, mean, or max) and $\text{MSG}^{(l)}$ and $\text{UPDATE}^{(l)}$ are learnable functions at layer l . Different choices of these three components recover different GNN architectures. GCN (Section 3.1.2) uses a linear message, a degree-based weighted sum aggregator, and a linear–nonlinear update. GAT (Section 3.1.3) introduces learned attention weights on the aggregator, making each neighbour’s contribution data-dependent. GraphSAGE [42] uses concatenated mean–max aggregation with a feedforward update, optimised for inductive settings.

The message-passing view clarifies two points that matter for the remainder of this chapter. First, the inductive biases of a GNN are encoded in the choice of aggregator and the edge-weighting scheme, not in the overall template. Second, the number of layers controls the receptive field — a node’s representation after L layers aggregates information from its L -hop neighbourhood — which, for skeleton graphs with at most four hops between any two joints, places a natural upper bound on depth.

3.1.2 Graph Convolutional Networks (GCN)

The GCN architecture of [43] derives from a first-order approximation of spectral graph convolutions. Given adjacency matrix $A \in \mathbb{R}^{N \times N}$ and degree matrix $D = \text{diag}(\sum_j A_{ij})$, the symmetric normalised adjacency with self-loops is

$$\tilde{A} = \tilde{D}^{-1/2} (A + I) \tilde{D}^{-1/2}, \quad (3.3)$$

where $\tilde{D} = \text{diag}(\sum_j (A + I)_{ij})$ is the degree matrix of the self-loop-augmented graph. The GCN propagation rule is

$$H^{(l+1)} = \sigma \left(\tilde{A} H^{(l)} W^{(l)} \right), \quad (3.4)$$

where $H^{(l)} \in \mathbb{R}^{N \times d_l}$ stacks the node features row-wise, $W^{(l)} \in \mathbb{R}^{d_l \times d_{l+1}}$ is a learnable weight matrix, and $\sigma(\cdot)$ is a pointwise nonlinearity.

Two properties of Equation (3.4) motivate the move to attention-based aggregation in the pose-grouping setting. First, the edge weights are fixed: once $\tilde{A}$ is constructed from the graph topology, no learning signal modulates the contribution of individual neighbours. Two neighbours of the same target node contribute in proportion determined only by the degrees of the three nodes involved. Second, this uniform treatment is informationally inadequate for skeleton graphs where edge semantics vary sharply: the edge from the left shoulder to the left elbow encodes different information from the edge from the

left shoulder to the head, and a pose-grouping embedding should be free to weight them differently. Attention-based aggregation, introduced next, provides exactly this flexibility.

3.1.3 Graph Attention Networks (GAT and GATv2)

GAT [44] replaces the fixed normalisation of Equation (3.3) with a learned attention weight for each edge. For a pair (i, j) with $j \in \mathcal{N}(i)$, the unnormalised attention score is

$$e_{ij}^{\text{GAT}} = \text{LeakyReLU}\left(\mathbf{a}^\top [W\mathbf{h}_i \| W\mathbf{h}_j]\right), \quad (3.5)$$

where $W \in \mathbb{R}^{d' \times d}$ is a shared linear projection, $\mathbf{a} \in \mathbb{R}^{2d'}$ is a learned attention vector, and $[\cdot \| \cdot]$ denotes vector concatenation. The scores are normalised over the neighbourhood with a softmax,

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}(i)} \exp(e_{ik})}, \quad (3.6)$$

and the updated node representation is the attention-weighted sum of projected neighbour features,

$$\mathbf{h}'_i = \sigma\left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij} W\mathbf{h}_j\right). \quad (3.7)$$

Multi-head attention is recovered by repeating Equations (3.5)–(3.7) with independent parameters for each head and concatenating or averaging the outputs.

Equation (3.5) exposes a subtle limitation that was identified by [45] and has direct consequences for pose grouping. Because $\mathbf{a}^\top$ is applied to the concatenation $[W\mathbf{h}_i \| W\mathbf{h}_j]$, the attention score decomposes as $\mathbf{a}_1^\top W\mathbf{h}_i + \mathbf{a}_2^\top W\mathbf{h}_j$ where $\mathbf{a} = [\mathbf{a}_1; \mathbf{a}_2]$. The term $\mathbf{a}_1^\top W\mathbf{h}_i$ is independent of j , so the relative ordering of neighbours by attention score is the same for every query node — a property that [45] formalises as *static attention*. In other words, if node j_1 is ranked above node j_2 as a neighbour of i , then j_1 is also ranked above j_2 as a neighbour of every other node. For a pose-grouping network this is undesirable: the network should be able to judge that a given left shoulder is important to a right shoulder but not to a left ankle of a different person.

GATv2 resolves the limitation by reordering the operations inside the attention score,

$$e_{ij}^{\text{GATv2}} = \mathbf{a}^\top \text{LeakyReLU}(W[\mathbf{h}_i \| \mathbf{h}_j]). \quad (3.8)$$

With the nonlinearity now placed inside the query–key interaction, the attention score does not decompose into independent per-node terms, and the ranking of neighbours is allowed to depend on the query node. GATv2 is strictly more expressive than GAT at the same parameter count [45], and every architecture evaluated in this dissertation uses GATv2 as its attention primitive. The remaining

GCN and GAT references in Chapters 4 and 5 refer to baselines and background, not to the working network.

3.2 PRELIMINARIES — CLUSTERING AND LEARNED GROUPING

3.2.1 k -means and constrained k -means

Classical k -means objective and Lloyd’s algorithm. Constrained k -means (must-link / cannot-link) as the family that encodes type constraints. Why k -means appears throughout this dissertation as the canonical “simple” clustering head, and the basis for PGA scoring with GT K .

3.2.2 Deep embedded clustering (DEC)

Xie et al. (2016) DEC objective: Student’s t -distribution soft assignments q_{ij} , target distribution p_{ij} , KL divergence minimisation. Why DEC was a natural candidate (learn embedding and clustering jointly), and how the alternating refinement works.

3.2.3 Slot attention

Locatello et al. (2020) slot attention: iterative attention between a set of learned slots and input features, soft-binding each input to a slot via competitive softmax. Why slot attention suggested itself for pose grouping (slots as person identities) and what the iteration assumes.

3.2.4 Graph partitioning (MinCut, NormCut)

Graph-cut objectives: minimum cut, normalised cut (Shi & Malik). Relaxation to continuous assignment matrices and the eigenvalue connection. How graph partitioning was applied as a learned grouping head.

3.2.5 Deep modularity networks (DMoN)

Tsitsulin et al. (2023) DMoN objective: Newman modularity as a soft, differentiable community-detection loss. The modularity matrix $B = A - dd^\top / 2m$ and the collapse-prevention regulariser. Why DMoN was the natural follow-up after the three learned heads failed, and what assumption it makes about the relationship between graph structure and cluster membership (this will matter in Chapter 5).

3.3 PROBLEM FORMULATION

Formalises the task. Input: set of N keypoints with types t_i and 2D positions $\mathbf{p}_i$. Output: a partition into K groups (one per person) with the type constraint. Evaluation metric: Pose Grouping Accuracy (PGA), defined as the Hungarian-matched joint assignment accuracy between predicted and GT groups.

3.4 DATASET PIPELINE

The synthetic data generator (controlled K , perfect GT, procedural skeletons), the COCO keypoints dataset (val2017 used throughout), and the train/val splits. Depth channel handling (retained or removed) and why that decision matters for COCO transfer.

3.5 BASELINE: STANDARD GAT + CONTRASTIVE + KNN

The baseline architecture: stacked GATv2 layers on a fully-connected graph over detected keypoints, contrastive loss pulling same-person keypoints together and pushing different-person keypoints apart, kNN clustering at test time. The depth-input variant and the no-depth variant. This is the starting point of the narrative.

3.6 LEARNED GROUPING HEADS

3.6.1 DEC applied to pose grouping

How DEC was instantiated on top of the GAT embeddings: initialisation strategy, cluster-count handling, training objective combined with contrastive loss.

3.6.2 Slot attention applied to pose grouping

Slot attention head on top of the GAT embeddings, with slots as person identities. Number of slots, initialisation, iteration count, and the binding-to-assignment readout.

3.6.3 Graph partitioning applied to pose grouping

Graph partitioning head on top of the GAT embeddings. Edge-weight construction, normalised-cut objective, and how the soft partition is read out into hard assignments.

3.7 SA-DMON — SKELETON-AWARE DEEP MODULARITY NETWORKS

A DMoN-based learned grouping head adapted for skeletons. v1 (learnable σ , clamped σ , median σ , spectral-term variants) and v2 (feature-only adjacency, entropy-only regularisation, full formulation). Ten experiments in total, all tried to fix the modularity misalignment in a different way.

3.8 PG-GAT — POSE GROUPING GRAPH ATTENTION NETWORK

3.8.1 Design intuition

Reframes the investigation: every learned head failed to beat kNN on the same embeddings, so the bottleneck is the embedding. PG-GAT strengthens the embedding while keeping the grouping head simple (kNN).

3.8.2 Type-pair attention modulation

Per-type-pair learnable bias in the GATv2 attention score, so that different joint-type relationships (e.g. left-shoulder-to-right-shoulder vs left-shoulder-to-head) can be weighted independently.

3.8.3 Skeleton-aware positional encoding

A positional encoding scheme that encodes relative skeleton position (not absolute image position), invariant to translation and scale.

3.8.4 Type-repulsion loss term

An auxiliary contrastive term that explicitly pushes same-type keypoints of different people apart (the difficult case: two left shoulders from different people must not collapse).

3.8.5 Training procedure

Two-stage training: synthetic pretraining (contrastive loss, all three modifications active), then COCO fine-tuning on detected keypoints (20 epochs). Hyperparameters, schedule, optimiser.

3.9 HYPERBOLIC PG-GAT

A variant that embeds the output into a Lorentz-model hyperbolic space. Euclidean internal attention, Lorentz output projection, logmap for clustering. Motivation: skeletons are trees, and trees embed more faithfully in hyperbolic space with fewer dimensions. The variant tests whether this translates into a quality or efficiency advantage for pose grouping.

3.10 TRIGAT — TRIPLET-BASED PRIMITIVE

A variant where the graph nodes are skeleton triplets (A–B–C chains, 19 per person) rather than individual joints. Pivot-angle features are rotation-invariant articulation descriptors. Triplets are embedded, clustered, and joint labels are voted from triplet-cluster assignments. Included to test whether a higher-order primitive improves the embedding.

3.11 INTEGRATION WITH A REALISTIC DETECTOR

How PG-GAT plugs into a HigherHRNet detection pipeline without joint training. Keypoint detection confidence filtering, the K-estimation head (brief treatment — the full version belongs to the PEng follow-up), end-to-end PGA computation.

3.12 EVALUATION PROTOCOL AND METRICS

PGA definition with Hungarian matching. Synthetic evaluation (100-image virtual test set, perfect GT). COCO evaluation zero-shot and after COCO fine-tuning. End-to-end evaluation using HigherHRNet

detections. What kNN-PGA measures, what COP-KM-PGA adds, and which numbers are the native MEng headlines (kNN) versus what is mentioned with a pointer to the PEng paper (type-constrained clustering).

CHAPTER 4 EXPERIMENTAL RESULTS

This chapter reports experimental results for every method presented in Chapter 3, in the order of the investigation. Each section follows the lab-book pattern (setup, table of numbers, interpretation). Failed approaches (learned heads, SA-DMoN, TriGAT, visual-feature augmentation) are reported with their numbers *and* explanations of why they failed. Honest caveats are flagged where the numbers are close calls.

4.1 BASELINE: STANDARD GAT

Results for the standard GAT baseline with and without the depth channel. The key observation: removing the MiDaS depth channel *improved* COCO transfer by roughly four points (MiDaS has a synthetic-to-real domain gap). Numbers: 0.994 synth kNN / 0.838 COCO kNN with depth; 0.901 / 0.879 without.

4.2 FAILED LEARNED GROUPING HEADS

4.2.1 DEC

DEC plateaued at 0.654 synth and did not transfer meaningfully to COCO. The Student's- t target distribution concentrates cluster mass but does not separate identities in the pose setting.

4.2.2 Slot attention

Slot attention underperformed kNN on the same embeddings (0.858 synth / 0.788 COCO with the depth baseline). The competitive softmax does not encode the type constraint.

4.2.3 Graph partitioning

Graph partitioning was closest to beating kNN on synthetic (0.954) but collapsed on COCO (0.752). Highly sensitive to edge-weight construction, does not transfer.

4.2.4 Cross-cutting finding

Every learned head failed to beat plain kNN on the same embeddings. This was the first evidence that the lever is the embedding, not the grouping head.

4.3 SA-DMON — TEN EXPERIMENTS

Full results table for all ten SA-DMoN variants (v1 with learnable / clamped / median σ , v1 with λ sweeps, v1 without spectral term, v2 with feature-adjacency variants, v2 with entropy-type-only regularisation, v2 full). Every variant plateaus in the 0.77–0.81 synthetic range. Root cause analysis is deferred to Chapter 5, but the empirical picture is: no SA-DMoN formulation closes the gap to kNN.

4.4 PG-GAT ABLATIONS

Four-row ablation: type-pair-only, pos-enc-only, repulsion-only, full. Each modification contributes independently; the full variant is the sum. Headline synthetic number: 0.910 synth kNN / 0.882 COCO kNN on the synthetic-only training regime. Table with explicit Δ against the standard-GAT-no-depth baseline.

4.5 COCO FINE-TUNING

PG-GAT fine-tuned on COCO keypoints (20 epochs). COCO kNN PGA rises from 0.882 to 0.957. Synthetic PGA rises from 0.910 to 0.928 (the fine-tune is not destructive to synthetic performance). This is the single largest gain of any experiment in the dissertation.

4.6 HYPERBOLIC PG-GAT

Two-part result. Pre-fine-tune: hyperbolic at $d = 32$ matches or slightly beats Euclidean at $d = 128$ (0.918 synth / 0.899 COCO kNN, vs Euclidean 0.910 / 0.882) — consistent with trees needing fewer dimensions in hyperbolic space. Post-fine-tune: the advantage disappears. Hyperbolic $d = 32$ COCO-fine-tuned reaches 0.943 kNN, still below Euclidean $d = 128$ COCO-fine-tuned (0.957). Reported honestly as an *efficiency trade* (comparable quality at $4\times$ smaller embedding), not a quality win.

4.7 TRIGAT — NEGATIVE RESULT

Triplet-based primitive underperforms joint-level PG-GAT on every comparison: 0.898 synth / 0.855 COCO with kNN, versus 0.910 / 0.882 for joint-level. Root cause: triplet clusters must be voted back to joint labels, and small triplet-level errors propagate to multiple joint votes. Documented as a negative result.

4.8 END-TO-END INTEGRATION WITH HIGHERHRNET

Replace the associative-embedding tag head in HigherHRNet with PG-GAT. End-to-end PGA using HigherHRNet-W32-512 for detection: PG-GAT + kNN = 0.942, PG-GAT + COP-Kmeans = 0.955 (COP-KM detail deferred to the PEng paper), HigherHRNet native AE grouping = 0.985. The gap to 0.985 comes from the tight coupling of the tag head to its backbone features, not from embedding-quality issues.

4.9 MODEL SIZE AND MODULARITY

Parameter counts: PG-GAT 170,560; K-head 18,625; HigherHRNet-W32 full model 28,645,331; HigherHRNet AE tag head alone ~561. Correction to the earlier “168× smaller” claim: PG-GAT is *larger* than the pure AE tag head (170K vs 561). The contribution is modularity — PG-GAT plugs into any detector’s keypoint output without joint training; the AE head cannot.

4.10 VISUAL-FEATURE AUGMENTATION (PG-GAT V2) — NEGATIVE RESULT

Added MobileNetV2 features sampled at each keypoint location as additional node input. Result: 0.968 COCO kNN / 0.964 E2E — a 0.8-point E2E improvement at the cost of 2.2M extra parameters. Total pipeline (30.8M) is larger than HigherHRNet (28.6M) while still scoring lower (0.964 vs 0.985). Reported as a negative result; the Paper 3 idea was dropped on this basis.

CHAPTER 5 DISCUSSION

This chapter interprets the results of Chapter 4, answers the research questions from Chapter 1, and examines the limitations of the work. The central question of the chapter is *why* the investigated methods succeeded or failed, not merely what the numbers were.

5.1 WHY THE LEARNED GROUPING HEADS FAILED

Discusses the shared failure mode of DEC, slot attention, and graph partitioning: each added a learned assignment mechanism on top of embeddings that were already close to separable under simple kNN. When the embedding is the bottleneck, stacking a learned head reshapes the assignment space but does not improve separation, and can actively hurt transfer (the learned head overfits to synthetic structure).

5.2 WHY SA-DMON FAILED

Root-cause analysis of the ten SA-DMoN experiments. The modularity objective is fundamentally misaligned with pose grouping: the kNN graph encodes *spatial* proximity, but the target communities are defined by *identity*. The spatial null model either cancels with the spatial adjacency or is miscalibrated, and no null-model reformulation (learnable / clamped / median σ , feature-only adjacency, entropy regularisation) repaired the misalignment. This is a structural, not a tuning, problem.

5.3 WHY PG-GAT WORKS

Discusses what each of the three modifications buys, drawing on the ablation numbers from Chapter 4. Type-pair attention lets the network weight joint-type relationships independently. The skeleton-aware positional encoding provides translation/scale invariance that the contrastive loss alone does not enforce. The type-repulsion term addresses the hardest pairs (same-type, different-person keypoints that look locally similar). Together they make the embedding separable enough that kNN works.

5.4 INTERPRETING THE HYPERBOLIC RESULT

Interprets the dim-32 hyperbolic advantage and its disappearance after fine-tuning. Pre-fine-tune, hyperbolic geometry helps because the model is relying on the geometric prior (trees embed naturally in hyperbolic space). After fine-tuning, the Euclidean model learns a sufficient representation from data, and the geometric prior becomes a constraint rather than a help. The result is a quality-matched efficiency trade, not a quality win.

5.5 WHY TRIGAT UNDERPERFORMED

The triplet primitive is semantically richer per-token but introduces a voting step that propagates errors from clusters to joint labels. A single mis-voted triplet cascades across the three joint votes that triplet contributes to, and the cascaded effect dominates the gap between TriGAT and joint-level PG-GAT in the results of Chapter 4.

5.6 COMPARISON WITH HIGHERHRNET AE — THE MODULARITY TRADE-OFF

Discusses the 0.955 vs 0.985 gap. HigherHRNet’s AE head is tiny (~ 561 parameters) but inseparable from its backbone; it gets quality from the backbone features being trained jointly with the tag. PG-GAT is larger (170K parameters) but detector-agnostic — a capability the AE head cannot offer. The correct comparison is not “who wins on COCO” but “what does each buy.” This framing matters for the modularity contribution claim.

5.7 ANSWERS TO THE RESEARCH QUESTIONS

Explicit answers to RQ1, RQ2, RQ3 from Chapter 1, grounded in the numbers from Chapter 4.

5.8 LIMITATIONS AND THREATS TO VALIDITY

Known limitations: reliance on detector-supplied K (K -estimation is outside the MEng core); synthetic-to-real domain gap mitigated but not eliminated; only one backbone evaluated end-to-end; PGA is Hungarian-matched joint accuracy, not mAP, so comparison to COCO-standard HPE numbers is indirect; COCO val2017 is the only real-image benchmark reported.

CHAPTER 6 CONCLUSION

This chapter summarises the contributions of the dissertation, consolidates the answers to the research questions, and indicates directions for future work. The PEng follow-up is the main pointer forward, but several MEng-local directions are also noted.

6.1 SUMMARY OF CONTRIBUTIONS

Restates the five numbered contributions from Section 1.6, each now grounded in the numbers from Chapter 4. PG-GAT is the main technical contribution. The empirical finding that the embedding dominates the grouping head is the main scientific contribution. The COCO fine-tuning protocol, the hyperbolic variant, and the TriGAT negative result are supporting contributions.

6.2 ANSWERS TO THE RESEARCH QUESTIONS

Condensed one-paragraph answers to RQ1 (yes — PG-GAT + kNN reaches 0.957 COCO), RQ2 (all three PG-GAT modifications contribute independently; COCO fine-tuning adds the largest single improvement), and RQ3 (PG-GAT is 0.03 PGA below HigherHRNet AE end-to-end, but is detector-agnostic — the trade-off is modularity for a small quality cost).

6.3 FUTURE WORK

Three directions. (1) The PEng follow-up re-frames grouping as structured assignment and develops SCOT, SCOT-KI, and THA — closing on HigherHRNet’s 0.985 with algorithmic rather than embedding changes. (2) A fully hyperbolic PG-GAT architecture (hyperbolic attention, not only hyperbolic output) would test the geometric-prior hypothesis more cleanly than the partial variant evaluated here. (3) End-to-end joint training with a detector would give up modularity to regain the coupling advantage AE exploits, and would quantify how much of the 0.03 gap is fundamental versus avoidable.

REFERENCES

- [1] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, “Microsoft COCO: Common objects in context,” in *European Conference on Computer Vision (ECCV)*, 2014.
- [2] Y. Chen, Y. Tian, and M. He, “Monocular human pose estimation: A survey of deep learning-based methods,” *Computer Vision and Image Understanding*, vol. 192, p. 102897, 2020.
- [3] C. Zheng, W. Wu, C. Chen, T. Yang, S. Zhu, J. Shen, N. Kehtarnavaz, and M. Shah, “Deep learning-based human pose estimation: A survey,” *ACM Computing Surveys*, vol. 56, no. 1, pp. 1–37, 2023.
- [4] M. Andriluka, L. Pishchulin, P. Gehler, and B. Schiele, “2D human pose estimation: New benchmark and state of the art analysis,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2014, pp. 3686–3693.
- [5] A. Toshev and C. Szegedy, “DeepPose: Human pose estimation via deep neural networks,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2014, pp. 1653–1660.
- [6] S.-E. Wei, V. Ramakrishna, T. Kanade, and Y. Sheikh, “Convolutional pose machines,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 4724–4732.

- [7] A. Newell, K. Yang, and J. Deng, “Stacked hourglass networks for human pose estimation,” in *European Conference on Computer Vision*. Springer, 2016, pp. 483–499.
- [8] B. Xiao, H. Wu, and Y. Wei, “Simple baselines for human pose estimation and tracking,” in *European Conference on Computer Vision*, 2018, pp. 466–481.
- [9] K. Sun, B. Xiao, D. Liu, and J. Wang, “Deep high-resolution representation learning for human pose estimation,” in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
- [10] B. Cheng, B. Xiao, J. Wang, H. Shi, T. S. Huang, and L. Zhang, “HigherHRNet: Scale-aware representation learning for bottom-up human pose estimation,” in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020.
- [11] K. He, G. Gkioxari, P. Dollár, and R. Girshick, “Mask R-CNN,” in *Proceedings of the IEEE International Conference on Computer Vision*, 2017, pp. 2961–2969.
- [12] H.-S. Fang, S. Xie, Y.-W. Tai, and C. Lu, “RMPE: Regional multi-person pose estimation,” in *Proceedings of the IEEE International Conference on Computer Vision*, 2017, pp. 2334–2343.
- [13] Z. Cao, T. Simon, S.-E. Wei, and Y. Sheikh, “Realtime multi-person 2D pose estimation using part affinity fields,” in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [14] S. Kreiss, L. Bertoni, and A. Alahi, “PifPaf: Composite fields for human pose estimation,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019, pp. 11 977–11 986.
- [15] A. Newell, Z. Huang, and J. Deng, “Associative embedding: End-to-end learning for joint detection and grouping,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

- [16] Z. Geng, K. Sun, B. Xiao, Z. Zhang, and J. Wang, “Bottom-up human pose estimation via disentangled keypoint regression,” in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021.
- [17] G. Brasó, N. Kister, and L. Leal-Taixé, “The center of attention: Center-keypoint grouping via attention for multi-person pose estimation,” in *IEEE International Conference on Computer Vision (ICCV)*, 2021, pp. 11 853–11 863.
- [18] Y. Li, S. Zhang, Z. Wang, S. Yang, W. Yang, S.-T. Xia, and E. Zhou, “TokenPose: Learning keypoint tokens for human pose estimation,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021, pp. 11 313–11 322.
- [19] Y. Yuan, R. Fu, L. Huang, W. Lin, C. Zhang, X. Chen, and J. Wang, “HRFormer: High-resolution transformer for dense prediction,” in *Advances in Neural Information Processing Systems*, vol. 34, 2021, pp. 7281–7293.
- [20] Y. Xu, J. Zhang, Q. Zhang, and D. Tao, “ViTPose: Simple vision transformer baselines for human pose estimation,” in *Advances in Neural Information Processing Systems*, vol. 35, 2022, pp. 38 571–38 584.
- [21] ———, “ViTPose++: Vision transformer for generic body pose estimation,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 46, no. 2, pp. 1212–1230, 2023.
- [22] D. Shi, X. Wei, L. Li, Y. Ren, and W. Tan, “End-to-end multi-person pose estimation with transformers,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022, pp. 11 069–11 078.
- [23] Y. Xiao, K. Su, X. Wang, D. Yu, L. Jin, M. He, and Z. Yuan, “QueryPose: Sparse multi-person pose regression via spatial-aware part-level query,” in *Advances in Neural Information Processing Systems*, 2022.
- [24] J. Yang, A. Zeng, F. Li, S. Liu, R. Zhang, and L. Zhang, “Explicit box detection unifies end-to-end multi-person pose estimation,” in *International Conference on Learning Representations (ICLR)*,

2023.

- [25] H. Liu, Q. Chen, Z. Tan, J. Liu, J. Wang, X. Su, X. Li, K. Yao, J. Han, E. Ding *et al.*, “Group pose: A simple baseline for end-to-end multi-person pose estimation,” in *IEEE International Conference on Computer Vision (ICCV)*, 2023.
- [26] X. Sun, B. Xiao, F. Liu, and Y. Wei, “Integral human pose regression,” in *European Conference on Computer Vision*, 2018, pp. 529–545.
- [27] A. Nibali, Z. He, S. Morgan, and L. Prendergast, “Numerical coordinate regression with convolutional neural networks,” *arXiv preprint arXiv:1801.07372*, 2018, arXiv preprint, not peer-reviewed.
- [28] J. Li, S. Bian, A. Zeng, C. Wang, B. Pang, W. Liu, and C. Lu, “Human pose regression with residual log-likelihood estimation,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021, pp. 11 025–11 034.
- [29] Y. Li, S. Yang, P. Liu, S. Zhang, Y. Wang, Z. Wang, W. Yang, and S.-T. Xia, “SimCC: A simple coordinate classification perspective for human pose estimation,” in *European Conference on Computer Vision*, 2022, pp. 89–106.
- [30] T. Jiang, P. Lu, L. Zhang, N. Ma, R. Han, C. Lyu, Y. Li, and K. Chen, “RTMPose: Real-time multi-person pose estimation based on MMPose,” *arXiv preprint arXiv:2303.07399*, 2023, arXiv tech report from OpenMMLab.
- [31] P. Lu, T. Jiang, Y. Li, X. Li, K. Chen, and W. Yang, “RTMO: Towards high-performance one-stage real-time multi-person pose estimation,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024, pp. 1491–1500.
- [32] T. Jiang, X. Xie, and Y. Li, “RTMW: Real-time multi-person 2d and 3d whole-body pose estimation,” *arXiv preprint arXiv:2407.08634*, 2024, arXiv tech report from OpenMMLab.

- [33] R. Khirodkar, T. Bagautdinov, J. Martinez, S. Zhaoen, A. James, P. Selednik, S. Anderson, and S. Saito, “Sapiens: Foundation for human vision models,” in *European Conference on Computer Vision*, 2024.
- [34] J. Yang, A. Zeng, R. Zhang, and L. Zhang, “X-Pose: Detecting any keypoints,” in *European Conference on Computer Vision*, 2024.
- [35] Z. Geng, C. Wang, Y. Wei, Z. Liu, H. Li, and H. Hu, “Human pose as compositional tokens,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023, pp. 660–671.
- [36] J. Li, C. Wang, H. Zhu, Y. Mao, H.-S. Fang, and C. Lu, “CrowdPose: Efficient crowded scenes pose estimation and a new benchmark,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019, pp. 10 863–10 872.
- [37] S.-H. Zhang, R. Li, X. Dong, P. Rosin, Z. Cai, X. Han, D. Yang, H. Huang, and S.-M. Hu, “Pose2Seg: Detection free human instance segmentation,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019, pp. 889–898.
- [38] Y. Xiu, J. Li, H. Wang, Y. Fang, and C. Lu, “Pose flow: Efficient online pose tracking,” in *British Machine Vision Conference*, 2018.
- [39] G. Ning, J. Pei, and H. Huang, “LightTrack: A generic framework for online top-down human pose tracking,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*, 2020, pp. 1034–1035.
- [40] M. Andriluka, U. Iqbal, E. Insafutdinov, L. Pishchulin, A. Milan, J. Gall, and B. Schiele, “PoseTrack: A benchmark for human pose estimation and tracking,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 5167–5176.
- [41] J. Gilmer, S. S. Schoenholz, P. F. Riley, O. Vinyals, and G. E. Dahl, “Neural message passing for quantum chemistry,” in *International Conference on Machine Learning (ICML)*, 2017.

REFERENCES

- [42] W. L. Hamilton, R. Ying, and J. Leskovec, “Inductive representation learning on large graphs,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [43] T. N. Kipf and M. Welling, “Semi-supervised classification with graph convolutional networks,” in *International Conference on Learning Representations (ICLR)*, 2017.
- [44] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, “Graph attention networks,” in *International Conference on Learning Representations (ICLR)*, 2018.
- [45] S. Brody, U. Alon, and E. Yahav, “How attentive are graph attention networks?” in *International Conference on Learning Representations (ICLR)*, 2022.

APPENDIX A DERIVATION OF HIGHER ORDER MODELS

A.1 EXAMPLE HEADING

A.1.1 Subheading example

[illegible]

Filler text to test template. Filler text to test template. Filler text to test template. Filler text to test
template. Filler text to test template. Filler text to test template. Filler text to test template. Filler text
to test template. Filler text to test template. Filler text to test template. Filler text to test template. Filler
text to test template. Filler text to test template. Filler text to test template. Filler text to test template.
Filler text to test template. Filler text to test template. Filler text to test template. Filler text to test
template. Filler text to test template. Filler text to test template. Filler text to test template. Filler text

[illegible]