

Skeleton-Aware Graph Embeddings for Multi-Person Pose Grouping

by

Dean Spencer Mills

Submitted in partial fulfillment of the requirements for the degree
Master of Engineering (Computer Engineering)

in the

Department of Electrical, Electronic and Computer Engineering
Faculty of Engineering, Built Environment and Information Technology

UNIVERSITY OF PRETORIA

January 2025

SUMMARY

Skeleton-Aware Graph Embeddings for Multi-Person Pose Grouping

by

Dean Spencer Mills

Supervisor: <to be completed>
Co-supervisor: <to be completed> [optional line]
Department: Electrical, Electronic and Computer Engineering
University: University of Pretoria
Degree: Master of Engineering (Computer Engineering)
Keywords: Keyword, keyword, keyword, keyword, keyword, keyword, keyword,
keyword, keyword, keyword, keyword, keyword

This is the EECE dissertation/thesis template for Latex, version 2025-11-24. Summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one summary paragraph one
summary paragraph one summary paragraph one summary paragraph one.

Summary paragraph two summary paragraph two summary paragraph two summary paragraph two
summary paragraph two summary paragraph two summary paragraph two summary paragraph two
summary paragraph two summary paragraph two summary paragraph two summary paragraph two
summary paragraph two summary paragraph two summary paragraph two summary paragraph two

summary paragraph two summary paragraph two summary paragraph two summary paragraph two
summary paragraph two summary paragraph two.

Summary paragraph three summary paragraph three summary paragraph three summary paragraph
three summary paragraph three summary paragraph three summary paragraph three summary paragraph
three summary paragraph three summary paragraph three summary paragraph three summary paragraph
three summary paragraph three summary paragraph three.

LIST OF ABBREVIATIONS

ANN	artificial neural network
EMI	electromagnetic interference
MCMC	Markov chain Monte Carlo
SVM	support vector machine

Note: Make sure the list of abbreviations is sorted alphabetically. This list is formatted as a table with no borders. Insert rows as needed.

LIST OF SYMBOLS [OPTIONAL]

I	current
R	resistance
V	voltage

Note: Make sure the list of symbols is sorted alphabetically, and/or according to the Greek alphabet. The list is formatted as a table with no borders. Insert rows as needed.

TABLE OF CONTENTS

CHAPTER 1	INTRODUCTION	1
1.1	CONTEXT	1
1.2	PROBLEM STATEMENT	1
1.3	RESEARCH GAP	1
1.4	RESEARCH OBJECTIVE AND RESEARCH QUESTIONS	1
1.5	APPROACH	2
1.6	RESEARCH CONTRIBUTIONS	4
1.7	RESEARCH OUTPUTS	4
1.8	OVERVIEW OF STUDY	5
1.9	DECLARATION REGARDING USE OF GENERATIVE AI	5
CHAPTER 2	LITERATURE REVIEW	6
2.1	MULTI-PERSON HUMAN POSE ESTIMATION	6
2.2	GROUPING IN BOTTOM-UP POSE ESTIMATION	6
2.3	GRAPH NEURAL NETWORKS FOR POSE	6
2.4	REPRESENTATION LEARNING FOR CLUSTERING	6
2.5	GEOMETRIC EMBEDDING SPACES	7
2.6	RESEARCH GAP	7
CHAPTER 3	METHODOLOGY	8
3.1	PRELIMINARIES — GRAPH NEURAL NETWORKS	8
3.1.1	The message-passing framework	8
3.1.2	Graph convolutional networks (GCN)	8
3.1.3	Graph attention networks (GAT and GATv2)	8
3.2	PRELIMINARIES — CLUSTERING AND LEARNED GROUPING	9
3.2.1	k -means and constrained k -means	9

3.2.2	Deep embedded clustering (DEC)	9
3.2.3	Slot attention	9
3.2.4	Graph partitioning (MinCut, NormCut)	9
3.2.5	Deep modularity networks (DMoN)	9
3.3	PROBLEM FORMULATION	9
3.4	DATASET PIPELINE	9
3.5	BASILINE: STANDARD GAT + CONTRASTIVE + KNN	10
3.6	LEARNED GROUPING HEADS	10
3.6.1	DEC applied to pose grouping	10
3.6.2	Slot attention applied to pose grouping	10
3.6.3	Graph partitioning applied to pose grouping	10
3.7	SA-DMON — SKELETON-AWARE DEEP MODULARITY NETWORKS	10
3.8	PG-GAT — POSE GROUPING GRAPH ATTENTION NETWORK	10
3.8.1	Design intuition	10
3.8.2	Type-pair attention modulation	10
3.8.3	Skeleton-aware positional encoding	11
3.8.4	Type-repulsion loss term	11
3.8.5	Training procedure	11
3.9	HYPERBOLIC PG-GAT	11
3.10	TRIGAT — TRIPLET-BASED PRIMITIVE	11
3.11	INTEGRATION WITH A REALISTIC DETECTOR	11
3.12	EVALUATION PROTOCOL AND METRICS	11
CHAPTER 4	EXPERIMENTAL RESULTS	12
4.1	BASILINE: STANDARD GAT	12
4.2	FAILED LEARNED GROUPING HEADS	12
4.2.1	DEC	12
4.2.2	Slot attention	12
4.2.3	Graph partitioning	12
4.2.4	Cross-cutting finding	13
4.3	SA-DMON — TEN EXPERIMENTS	13
4.4	PG-GAT ABLATIONS	13
4.5	COCO FINE-TUNING	13

4.6	HYPERBOLIC PG-GAT	13
4.7	TRIGAT — NEGATIVE RESULT	13
4.8	END-TO-END INTEGRATION WITH HIGHERHRNET	14
4.9	MODEL SIZE AND MODULARITY	14
4.10	VISUAL-FEATURE AUGMENTATION (PG-GAT V2) — NEGATIVE RESULT . .	14
CHAPTER 5	DISCUSSION	15
5.1	WHY THE LEARNED GROUPING HEADS FAILED	15
5.2	WHY SA-DMON FAILED	15
5.3	WHY PG-GAT WORKS	15
5.4	INTERPRETING THE HYPERBOLIC RESULT	16
5.5	WHY TRIGAT UNDERPERFORMED	16
5.6	COMPARISON WITH HIGHERHRNET AE — THE MODULARITY TRADE-OFF	16
5.7	ANSWERS TO THE RESEARCH QUESTIONS	16
5.8	LIMITATIONS AND THREATS TO VALIDITY	16
CHAPTER 6	CONCLUSION	17
6.1	SUMMARY OF CONTRIBUTIONS	17
6.2	ANSWERS TO THE RESEARCH QUESTIONS	17
6.3	FUTURE WORK	17
REFERENCES		18
APPENDIX A	DERIVATION OF HIGHER ORDER MODELS	19
A.1	EXAMPLE HEADING	19
A.1.1	Subheading example	19

LIST OF FIGURES

1.1 Synthetic dataset examples. (a) controlled single-person scene, (b) multi-person overlap stress case.	3
---	---

[THE LIST OF FIGURES IS OPTIONAL]

LIST OF TABLES

- 1.1 Grouping heads tested on standard-GAT-no-depth embeddings (synthetic-only training). The pattern motivates the investigation: the embedding, not the grouping head, is the decisive lever. PG-GAT keeps kNN as the head and strengthens the embedding. . 4

[THE LIST OF TABLES IS OPTIONAL]

CHAPTER 1 INTRODUCTION

This chapter motivates multi-person pose grouping as a separable problem within bottom-up human pose estimation (HPE), identifies the gap in modular detector-agnostic grouping methods, and sets out the research questions, approach, and contributions that the rest of the dissertation defends.

1.1 CONTEXT

Introduces the detection/grouping split in bottom-up HPE: modern detectors (HigherHRNet, HRNet) produce accurate per-joint heatmaps; assigning those joints to person identities is the remaining bottleneck. Summarises why joint-to-person assignment is intrinsically difficult (unknown K , type constraints, occlusion, scale) and why it deserves to be studied as a module in its own right.

1.2 PROBLEM STATEMENT

Formulates the task: given N detected keypoints with joint-type labels $t_i \in \{1, \dots, 17\}$ from an image containing an unknown number of people K , assign each keypoint to one of K person groups under the constraint that every person has at most one keypoint of each type.

1.3 RESEARCH GAP

Existing grouping methods (associative embeddings in HigherHRNet, part-affinity fields in OpenPose, HGG, CenterGroup) are tightly coupled to their host detectors and trained jointly with them. A *modular* grouping module that operates on any detector’s keypoint output is largely unstudied. This is the gap this work addresses.

1.4 RESEARCH OBJECTIVE AND RESEARCH QUESTIONS

The objective is to design and evaluate a detector-independent, GNN-based pose grouping module. The investigation addresses the following research questions. Bracketed short answers anticipate the findings from Chapter 4; full support is provided there.

- **RQ1 — The lever.** Is the bottleneck for clustering-based pose grouping the embedding, or the grouping algorithm? *[Answer: the embedding. Four learned grouping heads (DEC, slot attention, graph partitioning, SA-DMoN) fail to beat plain kNN on the same embeddings; strengthening the embedding (PG-GAT) then beats every prior configuration.]*
- **RQ2 — Inductive biases.** Which skeleton-aware modifications to a GATv2 embedding improve pose grouping, and by how much? *[Answer: three modifications — type-pair attention modulation, skeleton-aware positional encoding, and a type-repulsion loss term — each contribute independently; the full combination is the best configuration.]*
- **RQ3 — Detector independence.** Can a grouping module trained separately from the detector match an integrated end-to-end grouping mechanism (HigherHRNet’s associative embedding head) on a real-image benchmark? *[Answer: not fully. PG-GAT + kNN end-to-end is approximately 0.04 PGA below HigherHRNet AE on COCO, with detector-independence as the trade-off.]*
- **RQ4 — Embedding geometry.** Does the choice of embedding space (Euclidean vs hyperbolic) affect pose grouping quality? *[Answer: yes before COCO fine-tuning, where hyperbolic at $d = 32$ matches Euclidean at $d = 128$ (an efficiency edge). After fine-tuning the advantage disappears — reported as an efficiency trade, not a quality win.]*
- **RQ5 — Graph primitive.** Is a higher-order graph primitive (skeleton triplets) a better node for pose grouping than single joints? *[Answer: no. TriGAT underperforms joint-level PG-GAT on every comparison, with the voting step (triplet clusters to joint labels) as the dominant failure mode.]*
- **RQ6 — Fine-tuning.** How much does COCO fine-tuning lift a synthetically-trained embedding, and does it damage performance on the synthetic domain? *[Answer: +7.5 PGA on COCO (kNN), the single largest gain of any experiment in this work; synthetic performance is not damaged ($0.910 \rightarrow 0.928$).]*

1.5 APPROACH

Train a graph attention network with contrastive loss on detected keypoints; cluster at test time with kNN. The iteration order mirrors the investigation: standard GAT baseline $\rightarrow$ a series of learned grouping heads (DEC, slot attention, graph partitioning), all of which failed to beat kNN $\rightarrow$ SA-DMoN, a skeleton-aware modularity objective, which also failed $\rightarrow$ PG-GAT, which keeps the clustering head simple but strengthens the embedding with three skeleton-aware modifications. Evaluate on synthetic

data (controlled K , perfect GT) and COCO val2017 (zero-shot and fine-tuned). Close the loop by integrating with HigherHRNet for end-to-end validation.

1.6 RESEARCH CONTRIBUTIONS

The contributions of this work are:

1. **PG-GAT**, a skeleton-aware graph attention network for pose-keypoint embedding. Three modifications over standard GATv2: type-pair attention modulation, skeleton-aware positional encoding, and a type-repulsion term in the contrastive loss.
2. **Empirical evidence that the embedding, not the grouping head, is the decisive lever** for clustering-based pose grouping. Documented across four learned-head approaches (DEC, slot attention, graph partitioning, and ten SA-DMoN variants), none of which outperformed plain kNN on the same embeddings.
3. **A COCO fine-tuning protocol** that lifts PG-GAT to 0.957 PGA on COCO val2017 with kNN clustering, and end-to-end to 0.942 when paired with HigherHRNet detection.
4. **An alternative-embedding-space study** (hyperbolic PG-GAT in the Lorentz model): hyperbolic at $d = 32$ matches Euclidean at $d = 128$ pre-fine-tune, but the advantage disappears after COCO fine-tuning — documented as an efficiency trade, not a quality win.
5. **An alternative-primitive study** (TriGAT, triplet nodes): a negative result reported with the failure-mode analysis (voting step propagates errors).

1.7 RESEARCH OUTPUTS

The primary research output is a journal paper on PG-GAT, targeted at *Image and Vision Computing* (Q1, IF 5.34), with *Pattern Recognition* (Q1, IF 9.84) as a stretch target.

1.8 OVERVIEW OF STUDY

Chapter 2 reviews the literature on bottom-up human pose estimation, graph neural networks, and clustering methods relevant to pose grouping. Chapter 3 presents the mathematical preliminaries (GNN, GCN, GAT/GATv2, DEC, slot attention, graph partitioning, DMoN) and describes the methodology behind each method investigated, culminating in PG-GAT. Chapter 4 reports experimental results across the methods. Chapter 5 discusses what the results reveal about the problem. Chapter 6 concludes and indicates future work.

1.9 DECLARATION REGARDING USE OF GENERATIVE AI

Please declare if and how generative AI was used in this work. This must be done following the steps (adapted for dissertations/theses) outlined in 8.2.1.B.10 in the IEEE Publication Services and Products Board Operations Manual¹:

The use of content generated by artificial intelligence (AI) in ... [a dissertation/thesis] (including but not limited to text, figures, images, and code) shall be disclosed ... The AI system used shall be identified, and specific sections of the ... [dissertation/thesis] that use AI-generated content shall be identified and accompanied by a brief explanation regarding the level at which the AI system was used to generate the content.

While the use of AI systems for editing and grammar enhancement is common practice, its use in this role must also be disclosed.

¹<https://pspb.ieee.org/images/files/PSPB/opsmanual.pdf>

CHAPTER 2 LITERATURE REVIEW

This chapter situates multi-person pose grouping within the broader human pose estimation literature, reviews the existing families of grouping mechanisms, and surveys the graph neural network and representation-learning work that PG-GAT builds on. The review closes by identifying the specific gap that this dissertation addresses.

2.1 MULTI-PERSON HUMAN POSE ESTIMATION

Overview of top-down (detect-person-then-pose) versus bottom-up (detect-joints-then-group) pipelines. Why bottom-up is attractive (constant time in number of people, robust to overlapping detections) and why it reintroduces the grouping problem. Canonical detectors: Hourglass networks, HRNet, HigherHRNet.

2.2 GROUPING IN BOTTOM-UP POSE ESTIMATION

Review of the major grouping mechanisms in the bottom-up literature: part-affinity fields (OpenPose), associative embeddings (Newell et al.), PersonLab mid-range offsets, HGG (graph-based), CenterGroup. How each binds tightly to its host detector and why that limits modularity.

2.3 GRAPH NEURAL NETWORKS FOR POSE

Existing graph-based approaches in the pose literature: STGAT (skeleton-based action recognition), GAST-Net (single-person 3D pose from video), skeleton-aware GCN variants. How these differ from the multi-person keypoint grouping task — none of the existing “skeleton-aware” networks target grouping, which motivates both the PG-GAT contribution and the naming choice (avoiding collision with STGAT/GAST-Net descriptors).

2.4 REPRESENTATION LEARNING FOR CLUSTERING

Contrastive learning (InfoNCE, SimCLR), triplet losses, metric learning. Deep clustering lineage: DEC (Xie et al.), IDEC, DeepCluster, DMoN. What these methods assume about the feature space and how those assumptions interact with the pose-grouping task.

2.5 GEOMETRIC EMBEDDING SPACES

Euclidean versus non-Euclidean embedding spaces. Hyperbolic embeddings (Poincaré ball, Lorentz model) and their suitability for tree-structured data. Relevance to skeletons, which are trees. Sets up the motivation for the hyperbolic PG-GAT variant evaluated in Chapter 4.

2.6 RESEARCH GAP

Synthesises the review: modular grouping modules for bottom-up HPE are under-explored; existing graph-based pose networks target different tasks; the clustering-head literature has not been systematically applied to pose grouping; and the choice of embedding space (Euclidean vs hyperbolic) is untested in this domain. PG-GAT and its ablations address these gaps.

CHAPTER 3 METHODOLOGY

This chapter presents the mathematical preliminaries that every method in this dissertation builds on, the problem formulation and data pipeline, and the methodology for each approach investigated — from the standard GAT baseline, through the series of learned grouping heads that were tried and failed, through SA-DMoN, and finally to PG-GAT. Two variants (hyperbolic PG-GAT, TriGAT) and the integration with a realistic detector are described last.

3.1 PRELIMINARIES — GRAPH NEURAL NETWORKS

3.1.1 The message-passing framework

General formulation of graph neural networks as iterated message passing. Nodes, edges, neighbourhoods, aggregation, update. Why the message-passing view unifies GCN, GAT, and most variants used in this dissertation.

3.1.2 Graph convolutional networks (GCN)

Derivation from spectral graph theory to the Kipf–Welling simplification. The normalised-adjacency formulation $\tilde{A} = D^{-1/2}(A + I)D^{-1/2}$, the layer equation $H^{(l+1)} = \sigma(\tilde{A}H^{(l)}W^{(l)})$, and the assumptions it makes about edge uniformity. Why uniform edge weights are a limitation for skeletons where joint-pair relationships differ by type.

3.1.3 Graph attention networks (GAT and GATv2)

Original GAT (Veličković et al., 2018) attention: $e_{ij} = \text{LeakyReLU}(\mathbf{a}^\top [W\mathbf{h}_i \parallel W\mathbf{h}_j])$, followed by softmax over neighbours. The static-attention limitation (shared neighbour ranking across query nodes). GATv2 (Brody et al., 2022) fix: reorder to $e_{ij} = \mathbf{a}^\top \text{LeakyReLU}(W[\mathbf{h}_i \parallel \mathbf{h}_j])$, restoring dynamic attention. All experiments in this work use GATv2 as the attention primitive.

3.2 PRELIMINARIES — CLUSTERING AND LEARNED GROUPING

3.2.1 k -means and constrained k -means

Classical k -means objective and Lloyd’s algorithm. Constrained k -means (must-link / cannot-link) as the family that encodes type constraints. Why k -means appears throughout this dissertation as the canonical “simple” clustering head, and the basis for PGA scoring with GT K .

3.2.2 Deep embedded clustering (DEC)

Xie et al. (2016) DEC objective: Student’s t -distribution soft assignments q_{ij} , target distribution p_{ij} , KL divergence minimisation. Why DEC was a natural candidate (learn embedding and clustering jointly), and how the alternating refinement works.

3.2.3 Slot attention

Locatello et al. (2020) slot attention: iterative attention between a set of learned slots and input features, soft-binding each input to a slot via competitive softmax. Why slot attention suggested itself for pose grouping (slots as person identities) and what the iteration assumes.

3.2.4 Graph partitioning (MinCut, NormCut)

Graph-cut objectives: minimum cut, normalised cut (Shi & Malik). Relaxation to continuous assignment matrices and the eigenvalue connection. How graph partitioning was applied as a learned grouping head.

3.2.5 Deep modularity networks (DMoN)

Tsitsulin et al. (2023) DMoN objective: Newman modularity as a soft, differentiable community-detection loss. The modularity matrix $B = A - dd^\top / 2m$ and the collapse-prevention regulariser. Why DMoN was the natural follow-up after the three learned heads failed, and what assumption it makes about the relationship between graph structure and cluster membership (this will matter in Chapter 5).

3.3 PROBLEM FORMULATION

Formalises the task. Input: set of N keypoints with types t_i and 2D positions $\mathbf{p}_i$. Output: a partition into K groups (one per person) with the type constraint. Evaluation metric: Pose Grouping Accuracy (PGA), defined as the Hungarian-matched joint assignment accuracy between predicted and GT groups.

3.4 DATASET PIPELINE

The synthetic data generator (controlled K , perfect GT, procedural skeletons), the COCO keypoints dataset (val2017 used throughout), and the train/val splits. Depth channel handling (retained or

removed) and why that decision matters for COCO transfer.

3.5 BASELINE: STANDARD GAT + CONTRASTIVE + KNN

The baseline architecture: stacked GATv2 layers on a fully-connected graph over detected keypoints, contrastive loss pulling same-person keypoints together and pushing different-person keypoints apart, kNN clustering at test time. The depth-input variant and the no-depth variant. This is the starting point of the narrative.

3.6 LEARNED GROUPING HEADS

3.6.1 DEC applied to pose grouping

How DEC was instantiated on top of the GAT embeddings: initialisation strategy, cluster-count handling, training objective combined with contrastive loss.

3.6.2 Slot attention applied to pose grouping

Slot attention head on top of the GAT embeddings, with slots as person identities. Number of slots, initialisation, iteration count, and the binding-to-assignment readout.

3.6.3 Graph partitioning applied to pose grouping

Graph partitioning head on top of the GAT embeddings. Edge-weight construction, normalised-cut objective, and how the soft partition is read out into hard assignments.

3.7 SA-DMON — SKELETON-AWARE DEEP MODULARITY NETWORKS

A DMoN-based learned grouping head adapted for skeletons. v1 (learnable σ , clamped σ , median σ , spectral-term variants) and v2 (feature-only adjacency, entropy-only regularisation, full formulation). Ten experiments in total, all tried to fix the modularity misalignment in a different way.

3.8 PG-GAT — POSE GROUPING GRAPH ATTENTION NETWORK

3.8.1 Design intuition

Reframes the investigation: every learned head failed to beat kNN on the same embeddings, so the bottleneck is the embedding. PG-GAT strengthens the embedding while keeping the grouping head simple (kNN).

3.8.2 Type-pair attention modulation

Per-type-pair learnable bias in the GATv2 attention score, so that different joint-type relationships (e.g. left-shoulder-to-right-shoulder vs left-shoulder-to-head) can be weighted independently.

3.8.3 Skeleton-aware positional encoding

A positional encoding scheme that encodes relative skeleton position (not absolute image position), invariant to translation and scale.

3.8.4 Type-repulsion loss term

An auxiliary contrastive term that explicitly pushes same-type keypoints of different people apart (the difficult case: two left shoulders from different people must not collapse).

3.8.5 Training procedure

Two-stage training: synthetic pretraining (contrastive loss, all three modifications active), then COCO fine-tuning on detected keypoints (20 epochs). Hyperparameters, schedule, optimiser.

3.9 HYPERBOLIC PG-GAT

A variant that embeds the output into a Lorentz-model hyperbolic space. Euclidean internal attention, Lorentz output projection, logmap for clustering. Motivation: skeletons are trees, and trees embed more faithfully in hyperbolic space with fewer dimensions. The variant tests whether this translates into a quality or efficiency advantage for pose grouping.

3.10 TRIGAT — TRIPLET-BASED PRIMITIVE

A variant where the graph nodes are skeleton triplets (A–B–C chains, 19 per person) rather than individual joints. Pivot-angle features are rotation-invariant articulation descriptors. Triplets are embedded, clustered, and joint labels are voted from triplet-cluster assignments. Included to test whether a higher-order primitive improves the embedding.

3.11 INTEGRATION WITH A REALISTIC DETECTOR

How PG-GAT plugs into a HigherHRNet detection pipeline without joint training. Keypoint detection confidence filtering, the K-estimation head (brief treatment — the full version belongs to the PEng follow-up), end-to-end PGA computation.

3.12 EVALUATION PROTOCOL AND METRICS

PGA definition with Hungarian matching. Synthetic evaluation (100-image virtual test set, perfect GT). COCO evaluation zero-shot and after COCO fine-tuning. End-to-end evaluation using HigherHRNet detections. What kNN-PGA measures, what COP-KM-PGA adds, and which numbers are the native MEng headlines (kNN) versus what is mentioned with a pointer to the PEng paper (type-constrained clustering).

CHAPTER 4 EXPERIMENTAL RESULTS

This chapter reports experimental results for every method presented in Chapter 3, in the order of the investigation. Each section follows the lab-book pattern (setup, table of numbers, interpretation). Failed approaches (learned heads, SA-DMoN, TriGAT, visual-feature augmentation) are reported with their numbers *and* explanations of why they failed. Honest caveats are flagged where the numbers are close calls.

4.1 BASELINE: STANDARD GAT

Results for the standard GAT baseline with and without the depth channel. The key observation: removing the MiDaS depth channel *improved* COCO transfer by roughly four points (MiDaS has a synthetic-to-real domain gap). Numbers: 0.994 synth kNN / 0.838 COCO kNN with depth; 0.901 / 0.879 without.

4.2 FAILED LEARNED GROUPING HEADS

4.2.1 DEC

DEC plateaued at 0.654 synth and did not transfer meaningfully to COCO. The Student's- t target distribution concentrates cluster mass but does not separate identities in the pose setting.

4.2.2 Slot attention

Slot attention underperformed kNN on the same embeddings (0.858 synth / 0.788 COCO with the depth baseline). The competitive softmax does not encode the type constraint.

4.2.3 Graph partitioning

Graph partitioning was closest to beating kNN on synthetic (0.954) but collapsed on COCO (0.752). Highly sensitive to edge-weight construction, does not transfer.

4.2.4 Cross-cutting finding

Every learned head failed to beat plain kNN on the same embeddings. This was the first evidence that the lever is the embedding, not the grouping head.

4.3 SA-DMON — TEN EXPERIMENTS

Full results table for all ten SA-DMoN variants (v1 with learnable / clamped / median σ , v1 with λ sweeps, v1 without spectral term, v2 with feature-adjacency variants, v2 with entropy-type-only regularisation, v2 full). Every variant plateaus in the 0.77–0.81 synthetic range. Root cause analysis is deferred to Chapter 5, but the empirical picture is: no SA-DMoN formulation closes the gap to kNN.

4.4 PG-GAT ABLATIONS

Four-row ablation: type-pair-only, pos-enc-only, repulsion-only, full. Each modification contributes independently; the full variant is the sum. Headline synthetic number: 0.910 synth kNN / 0.882 COCO kNN on the synthetic-only training regime. Table with explicit Δ against the standard-GAT-no-depth baseline.

4.5 COCO FINE-TUNING

PG-GAT fine-tuned on COCO keypoints (20 epochs). COCO kNN PGA rises from 0.882 to 0.957. Synthetic PGA rises from 0.910 to 0.928 (the fine-tune is not destructive to synthetic performance). This is the single largest gain of any experiment in the dissertation.

4.6 HYPERBOLIC PG-GAT

Two-part result. Pre-fine-tune: hyperbolic at $d = 32$ matches or slightly beats Euclidean at $d = 128$ (0.918 synth / 0.899 COCO kNN, vs Euclidean 0.910 / 0.882) — consistent with trees needing fewer dimensions in hyperbolic space. Post-fine-tune: the advantage disappears. Hyperbolic $d = 32$ COCO-fine-tuned reaches 0.943 kNN, still below Euclidean $d = 128$ COCO-fine-tuned (0.957). Reported honestly as an *efficiency trade* (comparable quality at $4\times$ smaller embedding), not a quality win.

4.7 TRIGAT — NEGATIVE RESULT

Triplet-based primitive underperforms joint-level PG-GAT on every comparison: 0.898 synth / 0.855 COCO with kNN, versus 0.910 / 0.882 for joint-level. Root cause: triplet clusters must be voted back to joint labels, and small triplet-level errors propagate to multiple joint votes. Documented as a negative result.

4.8 END-TO-END INTEGRATION WITH HIGHERHRNET

Replace the associative-embedding tag head in HigherHRNet with PG-GAT. End-to-end PGA using HigherHRNet-W32-512 for detection: PG-GAT + kNN = 0.942, PG-GAT + COP-Kmeans = 0.955 (COP-KM detail deferred to the PEng paper), HigherHRNet native AE grouping = 0.985. The gap to 0.985 comes from the tight coupling of the tag head to its backbone features, not from embedding-quality issues.

4.9 MODEL SIZE AND MODULARITY

Parameter counts: PG-GAT 170,560; K-head 18,625; HigherHRNet-W32 full model 28,645,331; HigherHRNet AE tag head alone ~ 561 . Correction to the earlier “ $168\times$ smaller” claim: PG-GAT is *larger* than the pure AE tag head (170K vs 561). The contribution is modularity — PG-GAT plugs into any detector’s keypoint output without joint training; the AE head cannot.

4.10 VISUAL-FEATURE AUGMENTATION (PG-GAT V2) — NEGATIVE RESULT

Added MobileNetV2 features sampled at each keypoint location as additional node input. Result: 0.968 COCO kNN / 0.964 E2E — a 0.8-point E2E improvement at the cost of 2.2M extra parameters. Total pipeline (30.8M) is larger than HigherHRNet (28.6M) while still scoring lower (0.964 vs 0.985). Reported as a negative result; the Paper 3 idea was dropped on this basis.

CHAPTER 5 DISCUSSION

This chapter interprets the results of Chapter 4, answers the research questions from Chapter 1, and examines the limitations of the work. The central question of the chapter is *why* the investigated methods succeeded or failed, not merely what the numbers were.

5.1 WHY THE LEARNED GROUPING HEADS FAILED

Discusses the shared failure mode of DEC, slot attention, and graph partitioning: each added a learned assignment mechanism on top of embeddings that were already close to separable under simple kNN. When the embedding is the bottleneck, stacking a learned head reshapes the assignment space but does not improve separation, and can actively hurt transfer (the learned head overfits to synthetic structure).

5.2 WHY SA-DMON FAILED

Root-cause analysis of the ten SA-DMoN experiments. The modularity objective is fundamentally misaligned with pose grouping: the kNN graph encodes *spatial* proximity, but the target communities are defined by *identity*. The spatial null model either cancels with the spatial adjacency or is miscalibrated, and no null-model reformulation (learnable / clamped / median σ , feature-only adjacency, entropy regularisation) repaired the misalignment. This is a structural, not a tuning, problem.

5.3 WHY PG-GAT WORKS

Discusses what each of the three modifications buys, drawing on the ablation numbers from Chapter 4. Type-pair attention lets the network weight joint-type relationships independently. The skeleton-aware positional encoding provides translation/scale invariance that the contrastive loss alone does not enforce. The type-repulsion term addresses the hardest pairs (same-type, different-person keypoints that look locally similar). Together they make the embedding separable enough that kNN works.

5.4 INTERPRETING THE HYPERBOLIC RESULT

Interprets the dim-32 hyperbolic advantage and its disappearance after fine-tuning. Pre-fine-tune, hyperbolic geometry helps because the model is relying on the geometric prior (trees embed naturally in hyperbolic space). After fine-tuning, the Euclidean model learns a sufficient representation from data, and the geometric prior becomes a constraint rather than a help. The result is a quality-matched efficiency trade, not a quality win.

5.5 WHY TRIGAT UNDERPERFORMED

The triplet primitive is semantically richer per-token but introduces a voting step that propagates errors from clusters to joint labels. A single mis-voted triplet cascades across the three joint votes that triplet contributes to, and the cascaded effect dominates the gap between TriGAT and joint-level PG-GAT in the results of Chapter 4.

5.6 COMPARISON WITH HIGHERHRNET AE — THE MODULARITY TRADE-OFF

Discusses the 0.955 vs 0.985 gap. HigherHRNet’s AE head is tiny (~ 561 parameters) but inseparable from its backbone; it gets quality from the backbone features being trained jointly with the tag. PG-GAT is larger (170K parameters) but detector-agnostic — a capability the AE head cannot offer. The correct comparison is not “who wins on COCO” but “what does each buy.” This framing matters for the modularity contribution claim.

5.7 ANSWERS TO THE RESEARCH QUESTIONS

Explicit answers to RQ1, RQ2, RQ3 from Chapter 1, grounded in the numbers from Chapter 4.

5.8 LIMITATIONS AND THREATS TO VALIDITY

Known limitations: reliance on detector-supplied K (K -estimation is outside the MEng core); synthetic-to-real domain gap mitigated but not eliminated; only one backbone evaluated end-to-end; PGA is Hungarian-matched joint accuracy, not mAP, so comparison to COCO-standard HPE numbers is indirect; COCO val2017 is the only real-image benchmark reported.

CHAPTER 6 CONCLUSION

This chapter summarises the contributions of the dissertation, consolidates the answers to the research questions, and indicates directions for future work. The PEng follow-up is the main pointer forward, but several MEng-local directions are also noted.

6.1 SUMMARY OF CONTRIBUTIONS

Restates the five numbered contributions from Section 1.6, each now grounded in the numbers from Chapter 4. PG-GAT is the main technical contribution. The empirical finding that the embedding dominates the grouping head is the main scientific contribution. The COCO fine-tuning protocol, the hyperbolic variant, and the TriGAT negative result are supporting contributions.

6.2 ANSWERS TO THE RESEARCH QUESTIONS

Condensed one-paragraph answers to RQ1 (yes — PG-GAT + kNN reaches 0.957 COCO), RQ2 (all three PG-GAT modifications contribute independently; COCO fine-tuning adds the largest single improvement), and RQ3 (PG-GAT is 0.03 PGA below HigherHRNet AE end-to-end, but is detector-agnostic — the trade-off is modularity for a small quality cost).

6.3 FUTURE WORK

Three directions. (1) The PEng follow-up re-frames grouping as structured assignment and develops SCOT, SCOT-KI, and THA — closing on HigherHRNet’s 0.985 with algorithmic rather than embedding changes. (2) A fully hyperbolic PG-GAT architecture (hyperbolic attention, not only hyperbolic output) would test the geometric-prior hypothesis more cleanly than the partial variant evaluated here. (3) End-to-end joint training with a detector would give up modularity to regain the coupling advantage AE exploits, and would quantify how much of the 0.03 gap is fundamental versus avoidable.

REFERENCES

APPENDIX A DERIVATION OF HIGHER ORDER MODELS

A.1 EXAMPLE HEADING

A.1.1 Subheading example

[illegible]

Filler text to test template. Filler text to test template. Filler text to test template. Filler text to test
template. Filler text to test template. Filler text to test template. Filler text to test template. Filler text
to test template. Filler text to test template. Filler text to test template. Filler text to test template. Filler
text to test template. Filler text to test template. Filler text to test template. Filler text to test template.
Filler text to test template. Filler text to test template. Filler text to test template. Filler text to test
template. Filler text to test template. Filler text to test template. Filler text to test template. Filler text

[illegible]